

Medical Vision–Language Models Answer Through Broken Evidence

Hanqi Jiang, Junhao Chen, Mingyu Kang, Hyeokjae Kwon, Lifeng Chen, Weihang You, Yi Pan, Haozhen Gong, Caiwen Jiang, Zhengliang Liu, Wei Liu, Hui Ren, Quanzheng Li, Tianming Liu, and Xiang Li

Abstract—Medical vision–language models are evaluated on intact image–question pairs, but safe clinical use also requires a model to recognise when the evidence for an answer has failed. We formalise this requirement as an *evidence contract* and measure its violation as *silent failure*, a fluent non-refusal answer given under perturbed evidence. A 300-case suite authored by four board-certified radiologists shows the phenomenon in 2D across sixteen models. Capability and safe failure diverge, silent failure rises with clinical risk tier, and a radiologist blinded to construction leads the composite score by 14.1 points. Annotated labels limit such a suite to a few hundred items, and they cannot rule out pretraining overlap. We remove both limits by computing probe answers instead of annotating them. Whether a lesion would contact a critical structure after isotropic growth follows from one distance transform over an expert lesion mask and an automatic anatomy segmentation. This yields 9,484 verifiable probes over 587 CT volumes at no annotation cost, each asking about a state that never occurred. The silent-failure finding replicates and worsens. Dataset recognition measures 0.0–5.0% against a live positive control. The failure is localised: models perceive the annotation and perform the underlying numeric comparison, but they cannot estimate distance from the image, on anatomy or on a blank synthetic background, and their decision variable ranks the growth number in the question at AUROC up to 0.96 while ranking the computed label at 0.47–0.52. Two metrics validated in 2D do not survive the change of modality, and we report the corrections. The suite, harness and cached outputs are released.

Index Terms—Medical vision–language models, trustworthy AI, evaluation methodology, computed tomography, visual grounding, silent failure, counterfactual probing.

I. INTRODUCTION

MEDICAL vision–language models (VLMs) have moved from research prototypes into clinical imaging workflows, open and proprietary alike [1]–[5], in tasks from report assistance to decision support [6], [7]. As they approach deployment, accuracy on intact inputs is no longer a sufficient description of behaviour. A model must also behave safely when the answer-relevant region is obscured, when the

question embeds a premise the image contradicts, or when the requested quantity cannot be recovered from what was acquired.

Current systems often do not. Frontier VLMs reach top rank on medical benchmarks without receiving an image [8] and underperform young children on tasks requiring genuine perception [9]. In medicine this is hard to detect, because the literature is rich enough that a plausible answer can be assembled from question templates alone [10], [11]. A fabricated answer then reads like a grounded one and leaves no trace in an accuracy score.

Existing evaluation is not designed to detect this failure. Medical VQA datasets [12]–[14] and the systems trained on them [1], [6], [7] presume valid inputs, as does a concurrent audit of grounding pathologies [15]. Hallucination benchmarks [16]–[20] score unsupported content under otherwise valid inputs, measurement hallucinations are corrected after the fact in report generation [21], HEAL-MedVQA tests grounding through localisation [22], and selective prediction [23]–[25], abstention benchmarks [26]–[29] and risk-stratified evaluation [30], [31] treat refusal decision-theoretically. None of them targets the case in which giving an answer despite broken visual evidence is the unsafe act.

We formalise that case as an *evidence contract*. While the answer-relevant evidence is intact, the correct behaviour is to answer. Once the evidence has been removed, corrupted or contradicted, the correct behaviour is to recognise that the basis for an answer has failed. The measurable failure mode is *silent failure under perturbed evidence*, a non-refusal answer that accepts the perturbation (Fig. 1).

We first operationalise the contract in 2D with a 300-case suite supervised by four board-certified radiologists. Two annotate every case in parallel, a senior radiologist consolidates the manifest, and a fourth radiologist from a separate institution, blinded to construction, answers every probe as a human reference. An audit of sixteen models over 2,556 probes shows that capability and trustworthy behaviour do not reduce to one ranking, that silent failure rises with clinical risk tier, and that the blinded radiologist leads the composite by 14.1 points.

Two objections cannot be settled from inside that result. The first is scale. Four clinicians can author about 300 cases, and a reader may ask whether the finding survives an order of magnitude. The second is pretraining overlap. The suite draws on VQA-RAD, SLAKE, ROCO and public radiograph collections that almost certainly intersect the pretraining mixtures of every proprietary model audited. One can argue that the

This work has been submitted to the IEEE for possible publication. Copyright may be transferred without notice, after which this version may no longer be accessible.

H. Jiang, J. Chen, L. Chen, W. You, Y. Pan and T. Liu are with the University of Georgia, Athens, GA, USA. Z. Liu is with the School of Computing, University of Georgia, Athens, GA, USA. H. Jiang, Y. Pan, H. Ren, Q. Li and X. Li are with Harvard Medical School, Boston, MA, USA.


M. Kang is with Chungbuk National University, Cheongju, South Korea. H. Kwon is with Chungnam National University Hospital, Daejeon, South Korea. H. Gong is with the National University of Singapore, Singapore.

C. Jiang and W. Liu are with the Department of Radiation Oncology, Mayo Clinic, Phoenix, AZ, USA.

Corresponding author: X. Li (e-mail: XLI60@mgh.harvard.edu).

(a) 2D chest radiograph

Evidence perturbations expose unsupported answers



(b) 3D CT volume

Counterfactual growth tested against computed gold




Fig. 1. A fluent answer and a safe refusal look alike until the evidence is checked. (a) 2D. A frontal chest radiograph is asked “Is there a right apical pneumothorax?”, with gold “yes”. In the matched case the evidence is broken in two ways: the apex region is masked, and the question is edited to ask about a left apical pneumothorax. The silent failure reports that the left apical pneumothorax measures about 2.5 cm, fluently and against the evidence; the safe response states that the masked apex cannot be measured and that the finding is right apical, not left. (b) 3D. An axial abdominal CT with the liver lesion outlined in red, the counterfactual growth boundary dashed in yellow, and a magnified detail. The question asks whether the lesion would contact the highlighted vessel segment after 16 mm of isotropic growth; the surface gap is $d = 10.2$ mm, so the computed gold is “yes”. The silent failure answers “no” and contradicts the computed geometry; the correct answer states that the growth radius exceeds the minimum 3D gap. The CT is an illustrative case from the MSD hepatic-vessel task (CC BY-SA 4.0), outside the four audited tasks, and the model responses are illustrative.

perturbations postdate those corpora, which is sound, but it is an argument and not a measurement. Both objections share a root cause. The labels are annotated, so probes are expensive and tied to public images.

This article removes that root cause where it can be removed. In a volume, a useful class of question has an answer that follows from geometry. Given an expert lesion mask and an automatic anatomy segmentation, the question “if this lesion grew d mm in every direction, would it contact the aorta?” is decided by whether the surface gap is at most d . One Euclidean distance transform answers every such query for that lesion. Probe generation costs no clinician time, so the corpus is bounded by compute: 9,484 probes over 587 CT volumes of the Medical Segmentation Decathlon [32], against 2,556 probes over 300 cases. Because each probe asks about a state that never occurred, memorisation is excluded by construction.

In this paper we present a two-regime evaluation suite. It combines an annotation-backed 2D audit with a simulator-backed volumetric audit under one framework, metrics and composite, and it covers sixteen 2D and thirteen volumetric model configurations over 355,062 scored probes. Three results follow. First, the silent-failure finding replicates and worsens under computed labels, the 2D protocol’s text-side columns replicate at chance, and dataset recognition is measured directly instead of argued. Second, the failure is lo-

calised. Models perceive an outlined lesion, read the growth number in the question (their decision variable ranks it at AUROC up to 0.96), and perform the numeric comparison the task reduces to. They cannot estimate the distance between two outlined structures, on anatomy or on a blank synthetic background, and they do not apply the number inside a clinical sentence. Scaling from 3B to 72B parameters sharpens the reading of the number and leaves the label at chance. Third, two instruments validated in 2D do not transfer. The grounding contrast inverts, and likelihood scoring awards a perfect safety score to a constant refusal; both corrections apply to the 2D protocol as well. The suite, harness and cached outputs are publicly available.¹

II. MATERIALS AND METHODS

The evaluation comprises two suites under one framework (Table I): a 2D suite with clinician-authored labels and a volumetric suite with labels computed from geometry. Both are scored by the same metrics and composite, so the differences in Section III-C are differences in what the modality affords.

A. Evidence-Conditional Trustworthiness

Let f denote a medical VLM that maps an image I and a question q to a selected option $\hat{\ell} = f(I, q, \mathcal{O})$ from an option

¹Code: <https://github.com/hq0709/broken-evidence>. Data and cached model outputs: <https://huggingface.co/datasets/jhq0709/broken-evidence>.

TABLE I
THE TWO EVALUATION SUITES. BOTH ARE SCORED UNDER THE METRICS
OF SECTION II-D. CRT = CLINICAL RISK TIER.

	2D suite	Volumetric suite
Sources	VQA-RAD 120, SLAKE 60, ROCO 60, CXR 60	MSD [32]: liver 118, lung 63, pancreas 281, colon 125
Units	300 cases	587 CT volumes
Probes	2,556 MCQ	9,484 (8,476 growth-contact)
Counterfactuals	240 triplets	4,238 matched pairs
Text perturbations	PR, NEG, SDR by radiologists	same operators on the template
Options	five (A–E, E = refusal)	two (yes / no)
Label source	radiologist pipeline (3 stages)	distance transform
Annotation cost	four radiologists	none
Human reference	construction-blind radiologist	none (Sec. V)
CRT L1:L2:L3:L4:L5	71:31:118:43:37	assigned by target structure
Chance	20%	50% (balanced)

set $\mathcal{O}$. Let $\mathcal{P}$ be a set of evidence-perturbing operators acting on the textual or visual channel. A perturbation $\pi \in \mathcal{P}$ may rewrite the question (paraphrase, negation, specificity drop, knowledge-only rewrite, or hallucination trap) or replace the image with a region-of-interest variant (ROI-masked, ROI-only, laterality flip). For a perturbed probe whose constructed correct option is the refusal option ℓ_π^* , f exhibits *silent failure* on (I, q, π) when

$$\hat{\ell}_\pi(f) \neq \ell_\pi^*. \quad (1)$$

The refusal option is the safe action a clinician adjudicated for that probe (correct a false premise, flag a masked image, or decline). The definition separates the correctness of an answer from the prior question of whether any answer should be given.

B. The 2D Suite: Annotation-Backed Construction

The 2D suite contains 300 cases from VQA-RAD [12], SLAKE [13], ROCO [33], and chest radiograph collections from MIMIC-CXR [34] and CheXpert [35], stratified across five clinical risk tiers with a thirty-case minimum per tier; credentialed cases ship as reconstruction pointers. Each case carries the image, question, gold answer, a five-option set whose option E is the probe-specific safe action, a risk tier, a text-only-answerability flag, an ROI box and a laterality flag. All sources in both suites are public and de-identified; MIMIC-CXR and CheXpert are used under their data use agreements and MSD under CC BY-SA 4.0, no new patient data were collected, and no ethics approval was required.

Four board-certified radiologists build the suite in three stages (Fig. 2a). Two attending radiologists independently annotate each case and jointly author its text-side perturbations (paraphrase, negation, specificity drop, knowledge-only rewrite and two hallucination traps), while the three image-side variants are rendered from the ROI box. A senior radiologist adjudicates disagreements and rejects any perturbation that violates a construction invariant: a paraphrase must preserve the gold, a negation must invert it, a trap must contradict the image, and an ROI-masked variant must remove the answer-relevant region. Failed perturbations are dropped, not patched. A fourth radiologist from a separate institution, blinded to construction, answers every probe as the human reference.

C. The Volumetric Suite: Simulator-Backed Construction


The volumetric suite replaces annotation with computation. Every probe asks about a state that does not exist in the scan, so no report or language prior can supply the answer. Every answer follows from a deterministic rule over two masks, so it can be checked to the voxel.

1) *Materials*: We use the four Medical Segmentation Decathlon tasks [32] with lesion labels: liver, lung, pancreas and colon (587 volumes, Table I). Lesion instances L are the connected components of the expert lesion label with at least 20 voxels. Target structures A_k are 29 anatomical structures whose contact would change management (great vessels, heart, spinal cord, solid organs, bowel, lung lobes and selected vertebrae), segmented automatically with TotalSegmentator [36] in its 3 mm mode. The 3 mm mode reproduces the stored surface gaps to a median of 0.001 mm, whereas full-resolution segmentation deviates by a median of 0.54 mm and up to 1.96 mm, which would change the geometry the labels refer to.

2) *Computed counterfactual probes*: The primary family asks, “If this lesion grew by g mm in every direction, would it contact the A_k ?” Isotropic growth by g yields the set of voxels within g of L , so contact with A_k after growth holds if and only if the minimum lesion-to-target surface gap is at most g . One Euclidean distance transform [37] $D = \text{EDT}(\neg L)$, computed with the voxel spacing, therefore answers every growth query for that lesion: $\text{gap}(L, A_k) = \min_{v \in A_k} D(v)$. Algorithm 1 emits, for each (L, A_k) pair with a testable gap, one growth amount below the gap and one above it. The two members of a *matched pair* therefore have opposite answers by construction. Labels are balanced, chance is 50%, and a model that answers both members identically is not computing the geometry, which can be measured without reference to correctness (Section IV). A margin $m = 2$ mm removes lesions already in contact, and a 40 mm cap removes growth that is no longer clinically meaningful, so the corpus contains no gap below 2 mm (range 2.01–40.00 mm; Section V). Two secondary families use the same transform: *nearest-on-growth* (which of four comparably close structures is reached first) and *resection volume*. Risk tiers follow the target: aorta, heart and spinal cord L5, major organs L4, presence-of-finding L3, vertebrae L2, sternum and ribs L1.

3) *Rendering*: General-purpose VLMs receive the volume as a montage of three orthogonal slices through the volume centre, windowed to soft tissue (centre 40 HU, width 400 HU). The slices are resampled to isotropic pixels so that anisotropic slice spacing does not distort the vertical distances the model is asked to judge. This is the control arm of every input manipulation in Section III-B; native volumetric models receive the volume directly. The *identification control* holds the probe fixed and varies the input from `plain` (the montage above) through joint-visibility slices and red/cyan outlines of lesion and target to `identified` (both, plus a 10 mm scale bar drawn from the panel’s own spacing). An *input-richness* arm then holds annotation at that level and shows three planes, twenty-five contiguous axial slices, or the same planes cropped and magnified. Fig. 3 shows one probe under the four identification conditions.

(a) 2D · annotation-backed construction



(b) 3D · simulator-backed construction




Fig. 2. **Two constructions, one evaluation.** (a) 2D, annotation-backed. (1) Three hundred source cases from VQA-RAD (120), SLAKE (60), ROCO (60) and chest radiograph collections (60). (2) Two radiologists screen and annotate each case: they write image-dependent questions, remove text-answerable cases, and record the gold answer, A–E candidates, ROI and laterality, and a risk tier L1–L5. (3) Image and text perturbations are drafted as triplet candidates, controlled variants of the same case. (4) A senior radiologist finalises gold keys, checks construction invariants, resolves disagreements and drops invalid probes. (5) The bank holds vision-required MCQ probes with radiologist-adjudicated gold, risk labels and ROI annotations. (b) 3D, simulator-backed. (1) A CT volume with an expert lesion mask and an automatic anatomy segmentation. (2) One Euclidean distance transform per lesion, computed with physical voxel spacing, gives the surface gap d to each target structure. (3) A matched pair on the same lesion–target pair states one growth below the gap and one above it: for $d = 10.2$ mm, $g^- = 8$ mm gives no contact and $g^+ = 16$ mm gives contact. (4) A provenance gate checks gap, growth, spacing and masks and recomputes every label from the stored gap and growth. (5) The bank holds matched growth-contact probes with computed gold $y = 1[g \geq d]$, stored geometry and provenance. Both banks enter the shared evaluation of Fig. 4.

Algorithm 1 Computed counterfactual growth-contact probes for one lesion

Require: lesion mask L , target masks $\{A_k\}$, voxel spacing s , margin m , cap G , generator $\mathcal{U}$

- 1: $D \leftarrow \text{EDT}(\neg L; s)$ {one transform per lesion}
- 2: **for** each target A_k with $A_k \neq \emptyset$ **do**
- 3: $\text{gap} \leftarrow 0$ if $L \cap A_k \neq \emptyset$ else $\min_{v \in A_k} D(v)$
- 4: **if** $\text{gap} \leq m$ **or** $\text{gap} > G$ **then**
- 5: **continue** {vacuous or implausible}
- 6: **end if**
- 7: $g^- \leftarrow \max(1, \text{gap} - m - \mathcal{U}(0, 0.3 \text{ gap}))$
- 8: $g^+ \leftarrow \text{gap} + m + \mathcal{U}(0, 5)$
- 9: emit $(q(g^-), \text{no})$ and $(q(g^+), \text{yes})$ with provenance (gap, g)
- 10: **end for**
- 11: keep only complete pairs

4) *Verification:* Three checks verify the construction. The reference rule re-derived from each probe’s stored provenance reproduces the emitted label on 8,476 of 8,476 growth-contact probes. A regenerated segmentation cache reproduces the stored gaps to a median of 0.001 mm with zero answer flips. A render-parity test on a synthetic volume asserts that the `plain` arm is bit-exact against the published rendering path and that the annotated arms agree with it outside their annotation. The last check is a property of pixels, since two different renderings can both score 50%.

5) *Text-side perturbations:* The 2D suite perturbs the question while the image stays intact. The volumetric probe is a template, so the same three operators apply to every item and the gold stays computed: a *paraphrase* (PR) rewrites the question with one of three templates, chosen per matched pair so that both members share it, and preserves the gold; a *negation* (NEG) asks whether the grown lesion would *stay clear of* the target and inverts the gold; a *specificity drop* (SDR) removes the qualifier “in every direction” and preserves the gold, since isotropic growth is the default reading. All three run on the published `plain` rendering with the scoring path unchanged.

6) *Metric-channel control:* The distance sub-task asks how many millimetres separate the red-outlined lesion from the cyan-outlined structure. A model that cannot read a metric distance off any image would fail it for reasons unrelated to CT, so the same question, legend, options, scale bar, colours and scoring rule are asked on three backgrounds: the item’s real rendering; the real rendering with its annotations erased and two synthetic outlines drawn over the anatomy at a known surface gap; and three uniform-grey panels with the same synthetic outlines. The synthetic gaps follow the gap distribution of the 300 real items, and each gold is recomputed from the drawn masks with the same distance transform the corpus uses. The pixel scale of the synthetic bar is the panel’s own, so the three backgrounds differ only in what surrounds the outlines.




Fig. 3. **What the volumetric models are shown.** One growth-contact probe (colon, lesion 1, target liver; surface gap 3.97 mm, so growth of 1.0 mm gives “no” and 7.5 mm gives “yes”) rendered by the released harness. (a–d) The axial panel under the four identification conditions. The published `plain` input takes the three centre slices, on which the lesion the question names has no voxel, as for 42.0% of growth-matched probes; `bestslice` re-points the panels at slices on which both structures are visible; `overlay` keeps the centre slices and adds the legend sentence, so the prompt describes a red outline the image does not contain; `identified` combines joint-visibility slices, both outlines and a 10 mm bar drawn from each panel’s own spacing, and is byte-identical to the reader-study rendering and the input of the sub-task decomposition (Fig. 7). (e) The three panels of the `identified` montage as the model receives them, cropped to the body and shown at one physical scale. For display the outlines are drawn as contours of the resampled masks; the models receive the uncropped 1-px pixel outlines.

D. Metrics and Composite Score

Every metric is an exact-match rate over a probe family; no judge model is used. Fig. 4 summarises the shared evaluation.

Capability is Acc_{orig} on original probes and, in 2D, paraphrase (PR), negation (NEG) and specificity-drop (SDR) accuracy. *Evidence safety* is the silent-failure rate on traps, $\text{SFR}(f) = \Pr_{j \in \mathcal{J}_{\text{trap}}}[\hat{\ell}_j(f) \neq \ell_j^*]$. *Visual grounding* is the case-paired contrast $\text{VGR} = \text{Acc}_{\text{ROI-only}} - \text{Acc}_{\text{ROI-masked}}$, read alongside $\text{Acc}_{\text{ROI-masked}}$. *Language-prior accuracy* (LPA) is accuracy on knowledge-only probes answerable from text alone.

Because an L1 modality trap and an L5 don’t-miss trap should not weigh equally, the composite uses a harm-weighted rate,

$$\text{SFR}_w(f) = \frac{\sum_{k=1}^5 w_k \text{SFR}_{L_k}(f)}{\sum_{k=1}^5 w_k}, \quad (2)$$

with $(w_1, \dots, w_5) = (1, 2, 3, 5, 8)$ placing an L5 silent failure at eight times the weight of an L1 one, and three axes on $[0, 100]$: $\text{Cap} = \frac{1}{4}(\text{Acc}_{\text{orig}} + \text{PR} + \text{NEG} + \text{SDR})$, $\text{Safe} = 100 - \text{SFR}_w$, and $\text{Ground} = \frac{1}{2}(\text{clip}(\text{VGR} + 50, 0, 100) + \text{Acc}_{\text{ROI-masked}})$. The composite score (CS) is their harmonic

mean,

$$\text{CS}(f) = \frac{3 \text{ Cap Safe Ground}}{\text{Cap Safe} + \text{Cap Ground} + \text{Safe Ground}}, \quad (3)$$

which penalises the weakest axis (a 90/90/10 model scores ≈ 25 , not 63), so clean-input capability cannot offset unsafe trap behaviour or weak grounding. The headline SFR column is the unweighted rate; the Safe axis uses SFR_w .

1) *Scoring rule*: The 2D suite is scored by the single letter each model emits. The volumetric suite is scored by *likelihood over the option strings*, never by parsing free text, because native volumetric models answer in referring-expression templates whose wording embeds the answer words. Over options of unequal length the likelihood favours the string, not the content [38], so a refusal string can pin the argmax to one option for every item. On a balanced corpus this scores 50%, which an accuracy table cannot distinguish from chance. Two controls therefore accompany every volumetric result. *Content-free calibration* [39] subtracts the answer-string prior measured on a contentless question against a common baseline. *Greedy-decoding variants* (suffix `-gen`) parse a decoded answer, so that a scoring artefact and a genuine chance result can be told apart.




Fig. 4. **One evaluation for both suites.** Standard VQA asks whether an answer is correct; the evaluation asks whether the answer remains supported as the evidence changes. (a) Matched cases: a 2D chest radiograph with a five-option question and radiologist-adjudicated gold, and a 3D CT volume with a yes/no growth-contact question and computed gold with provenance, each posed with valid evidence and a well-posed premise. (b) Controlled evidence interventions. The visual evidence is full, ROI-masked, ROI-only or blank in both modalities; the textual premise is paraphrased, negated or replaced by a false premise; the modality-specific controls are a laterality flip in 2D and a matched growth pair, $g^- < d < g^+$, in 3D. (c) Paired response audit: the model answers the clean, visually edited and textually edited inputs, and each answer is parsed and matched against physician gold in 2D or computed gold in 3D, so that answer changes are tracked across the pair. (d) Each response is classed against the evidence available as a supported answer, a safe abstention or premise correction, expressed uncertainty, or a silent failure. (e) Capability (task accuracy), safety (silent-failure rate) and grounding (VGR) feed the risk-weighted composite, with diagnostic breakdowns for the text-side perturbations, sighted versus blind inputs, ROI integrity and the response-channel control.

2) *Statistical analysis:* Every metric is an exact-match rate, and no hypothesis test is used to declare a capability. Intervals are 95% percentile bootstrap intervals clustered by case in 2D and by volume in 3D, $B = 10,000$ unless stated; contrasts between arms are paired differences with clustered intervals on the same units. The identification control tests twelve paired contrasts on two endpoints, so about one nominal 95% interval is expected to exclude chance, and no single excluded interval is read as capability. A model that places a share m of its answers on one option of a balanced binary probe has accuracy confined to $[m - \frac{1}{2}, \frac{3}{2} - m]$ regardless of what it computes, so the modal share is reported beside every volumetric accuracy. Threshold-free ranking is measured by the AUROC of the decision variable $\log p_{\text{yes}} - \log p_{\text{no}}$ with the same clustered intervals. Robustness of the composite ordering is assessed by the Spearman correlation between CS rankings under alternative harm weights and grounding normalisations.

E. Audited Models and Protocol

Table II lists the 29 audited configurations with the two text-only baselines and the scale arm. 2D models are queried at $T = 0$ with a single-letter prompt, no chain of thought and no voting; two text-only baselines bound the language-prior contribution. The thirteen volumetric models span eight families and both input routes. The three native systems were chosen because they disagree in vision tower, backbone and volume-axis convention. Aria-25B-MoE was excluded because its checkpoint left 338 language-model tensors uninitialised, a

TABLE II
AUDITED MODEL CONFIGURATIONS. MONTAGE-FED MODELS RECEIVE THREE ORTHOGONAL SLICES; NATIVE MODELS RECEIVE THE VOLUME. ALL OPEN-WEIGHT MODELS LOAD THROUGH TRANSFORMERS; THE 2D PROPRIETARY ROWS ARE A DATED SNAPSHOT.

Suite	Family	Models	Input
2D	OpenAI	GPT-5.5, 5.4, 4o, 5.4-mini, 5.4-nano	API
	Anthropic	Claude Opus 4.7, Sonnet 4.6, Haiku 4.5	API
	Google	Gemini 3 Flash, 3.1 Flash-Lite	API
	Qwen	Qwen3.5-9B, Qwen3.5-397B-A17B	API
	Moonshot	Kimi-K2.5, K2.6	API
	medical, open text-only	LLaVA-Med-v1.5-7B, HuatuoGPT-Vision-7B, DeepSeek-V4-Flash, V4-Pro	open no image
3D	Qwen	Qwen2.5-VL 3B / 7B / 32B, Qwen3-VL-8B	montage
	InternVL	InternVL3 8B / 14B	montage
	other open	SmolVLM2-2.2B, LLaVA-OneVision-7B, Idefics3-8B, Pixtral-12B	montage
	native medical	M3D-LaMed-Phi3-4B, M3D-LaMed-Llama2-7B, Med3DVLM-7B	volume
	scale arm	Qwen2.5-VL-72B (identification control only)	montage

load the harness now refuses. InternVL3-78B and Llama-3.2-90B-Vision exceed two 80 GB cards and were not run, since quantising would change the claim under test.

1) *Response-channel control:* Before any volumetric score is interpreted, each model answers seven known-answer questions with no clinical content, such as “Is this a CT scan?”, on twenty volumes through the identical scoring path. A model that cannot do so has a pinned output channel, and its scores are not interpretable. The control is unnecessary in 2D, where

every model passes, and a prerequisite in 3D (Section III-B). It is run under both native and montage framings inside one library stack.

III. RESULTS

We report the 2D audit first, then the volumetric audit, then the two regimes against each other. Every number is recomputed from the released per-probe outputs; intervals are 95% bootstrap intervals clustered by case (2D) or by volume (3D), $B = 10,000$.

A. 2D Audit

Table III reports the sixteen vision-capable models and the radiologist reference on all seven metrics and the composite; Fig. 5 shows the three axes. Claude Opus 4.7 leads CS at 69.2 with the highest Capability, the lowest SFR and a positive VGR; Gemini 3.1 Flash-Lite is second at 66.0 on the strength of Safety and Grounding; GPT-4o falls to 44.1 because Safety is its bottleneck; Qwen3.5-9B is lowest at 31.0. The construction-blind radiologist scores 83.3 at 5.8% SFR, 14.1 points above the strongest model. Every audited model has LPA above 75%, so failures elsewhere are not knowledge deficits. The two text-only baselines, DeepSeek-V4-Flash and V4-Pro, score 4.0% and 1.7% on original probes and 99.3% on LPA, so original accuracy depends on the image. The ordering is stable. Case-clustered bootstrap intervals (500 resamples) put the leader at CS 69.2 [66.6, 72.1], against 66.0 [62.8, 68.2] for Gemini 3.1 Flash-Lite and 61.3 [58.2, 64.3] for GPT-5.5. Under nine alternative composites (linear, uniform, exponential and high-skew harm weights, and four grounding normalisations) the Spearman correlation with the default ranking is at least 0.97, the top five are unchanged, and the leader is the same.

Capability and safe behaviour do not move together. Gemini 3 Flash reaches 71.9% Overall at 37.2% SFR; GPT-4o 56.1% at 53.0%; Kimi-K2.5 47.3% at 60.0%. The two open-weight medical systems reach CS 44.7 (LLaVA-Med 7B) and 50.4 (HuatuoGPT-V 7B), roughly twenty points below the leader.

Silent failure rises with clinical risk tier for most models. GPT-4o answers 87.3% of L1 original probes correctly and fabricates on 68.2%, 75.6% and 68.9% of L3, L4 and L5 traps; GPT-5.4-mini shows 69.8% and 67.6% on L4 and L5. Claude Opus 4.7, the safest model, retains 30.9% L3 and 28.4% L5 trap SFR, against the radiologist’s 5.9% and 12.2%. The risk-weighted SFR_w ranges from 25.2 (Opus 4.7) to 76.1 (Qwen3.5-9B); the radiologist is at 9.3.

The paired ROI design separates two unsafe regimes. Gemini 3 Flash is vision-anchored but unsafe (VGR +41.1 pp, SFR 37.2%). It uses the ROI when present and fails to refuse when the ROI is removed. GPT-5.4-nano is prior-driven and unsafe (−27.1 pp, 51.7%), and it does not use ROI evidence consistently. The two regimes can be separated because VGR and trap SFR vary independently.

We sweep an isotropic Gaussian blur at $\sigma \in \{0, 2, 4, 8, 16, 32, 64\}$ px with a no-image control and define the language-takeover point L^* as the smallest σ at which the residual visual contribution falls below 20%

of the clean-image contribution. Across four representative models L^* varies by a factor of four (16–64 px) and does not track CS (Fig. 6). GPT-4o is the most visually anchored ($L^* = 64$). Claude Sonnet 4.6 has the largest visual contribution (65.9% $\rightarrow$ 9.1%) but falls back at $\sigma = 32$. Qwen3.5-397B has the smallest contribution (40.5 pp) and the earliest takeover ($L^* = 16$). HuatuoGPT-V 7B has the highest no-image floor (20.6%, four times that of GPT-4o).

B. Volumetric Audit

1) *Response-channel control*: Eight montage-fed models answer all 140 known-answer controls correctly at the 42.9% yes-rate the set implies (Table IV). SmolVLM2-2.2B fails one phrasing (121/140). Four models answer from a fixed direction instead of from the image. M3D-LaMed-Phi3-4B passes 107/140 (97% on yes-gold, 61% on no-gold), M3D-LaMed-Llama2-7B 80/140 (100% / 25%), Idefics3-8B 82/140 (3% / 98%), and Med3DVLM-7B 80/140 (0% / 100%). All three native volumetric systems and one general-purpose montage model fail the gate, and this set is not the one an architectural reading would predict. Run under both framings inside one library stack, native framing recovers M3D-LaMed-Phi3 by +29.3 points (74.3% native against 45.0% montage) and M3D-LaMed-Llama2 by +5.7. Med3DVLM answers “no” to all 140 controls under either framing.

2) *Headline*: The growth-matched subset is constructed so that a single threshold on the growth amount scores 50.8%, against 69.2% on the full corpus. On this subset all thirteen models lie between 48.7% and 51.0%, and every interval contains 50 (Table IV). Image contributions (sighted minus blind) run from −1.3 to +1.0 points. Seven of the thirteen models place at least 96% of their answers on one option, and every model answers at least 79.4% of matched pairs identically. The result does not depend on a threshold. The decision variable $\log p_{\text{yes}} - \log p_{\text{no}}$ ranks the computed label at AUROC 0.470–0.522 across the thirteen models (Fig. 9a), so no re-thresholding of any model’s scores recovers the label. Two independent measurements show that the volumetric pathway is not inert. Replacing the volume with zeros costs 4.3 to 10.4 accuracy points across four organs and two model families, and the image changes a large share of answers for five models (below).

3) *Text-side perturbations*: Under the three text-side operators the four clean-channel models stay at chance (Table V): paraphrase 47.8–51.4%, negation 49.5–50.0% and specificity drop 49.4–50.0%, every interval containing 50. Negation is ignored rather than mishandled. Every model answers “no” on 99.2–100% of negated probes, so a question and its negation receive the same answer; the share of negated probes answered consistently with the model’s own answer to the original is 0.0% for Qwen2.5-VL-7B, 27.6% for InternVL3-8B, 52.3% for Qwen2.5-VL-32B and 79.5% for Qwen3-VL-8B, the last only because its original answers were mostly “yes”. Under paraphrase the models keep their original answer on 54.4–100% of items. In 2D the same three columns span 29.7–80.7%, 15.6–83.1% and 25.0–79.9% and negation costs every model accuracy relative to the original; in 3D there is no

TABLE III

2D HEADLINE RESULTS (%; VGR IN PERCENTAGE POINTS). THE ITALIC ROW IS THE CONSTRUCTION-BLIND RADIOLOGIST, SHOWN AS A CALIBRATION ANCHOR RATHER THAN A LEADERBOARD ENTRY. ↓ LOWER IS SAFER. † TRAINED ON PUBMEDVISION, WHICH OVERLAPS VQA-RAD AND SLAKE.

Model	Capability		Evidence safety		Text robustness			Prior	Composite
	Overall	Original	SFR ↓	VGR	PR	NEG	SDR	LPA	CS
Radiologist (<i>ref.</i>)	92.4	95.3	5.8	+4.5	92.7	90.7	91.5	93.6	83.3
GPT-5.5	69.4	75.0	36.8	+12.0	74.3	68.9	59.8	99.6	61.3
GPT-5.4	65.2	68.0	28.8	+15.6	68.7	56.4	59.8	98.6	57.9
GPT-4o	56.1	59.3	53.0	+3.6	63.3	45.8	50.6	98.6	44.1
GPT-5.4-mini	54.7	55.3	49.7	−8.9	57.0	44.9	47.0	98.2	42.9
GPT-5.4-nano	38.8	31.7	51.7	−27.1	29.7	16.9	25.0	90.5	31.6
Claude Opus 4.7	76.5	78.7	22.0	+22.9	78.0	83.1	79.9	98.6	69.2
Claude Sonnet 4.6	61.2	63.3	41.3	+18.2	66.3	60.9	56.7	97.9	52.9
Claude Haiku 4.5	49.6	46.7	49.0	−9.4	51.0	32.4	42.1	97.5	39.9
Gemini 3 Flash	71.9	77.0	37.2	+41.1	80.7	75.6	76.2	98.9	63.7
Gemini 3.1 Flash-Lite	70.7	73.0	22.3	+49.5	75.7	63.1	70.7	98.9	66.0
Qwen3.5-9B	39.4	43.3	62.8	+12.5	48.0	28.4	32.3	84.5	31.0
Qwen3.5-397B-A17B	46.8	53.0	54.3	+24.0	53.0	33.3	42.7	92.2	39.9
Kimi-K2.5	47.3	49.7	60.0	+10.4	54.3	36.4	41.5	95.1	38.2
Kimi-K2.6	46.6	49.3	50.8	−2.1	47.7	36.0	37.8	95.4	38.8
LLaVA-Med-7B	44.8	43.3	43.0	+28.6	50.0	15.6	50.6	75.6	44.7
HuatuoGPT-V-7B†	55.5	55.3	31.5	+28.1	59.7	26.2	51.2	91.9	50.4




Fig. 5. **No model is strong on all three axes.** 2D audit summary: the three composite axes per model with the harmonic-mean composite score; models sort by CS; the radiologist reference is given in Table III.

accuracy to lose, and the operator that should invert the answer leaves it where it was.

4) *Identification control:* In the published condition the lesion the question names has no voxel on any slice shown for 42.0% of growth-matched probes, and both structures appear together on only 54.2%. Table VI holds the probe fixed and varies the input (Fig. 3). The eighteen cells, four conditions for four models and two for Qwen2.5-VL-72B, span 48.2–52.6%.

The two largest, Qwen3-VL-8B at 52.4 [50.3, 54.6] and Qwen2.5-VL-32B at 52.6 [50.7, 54.6] under identified, lie within three points of chance, which is the size of exclusion expected about once among twelve paired contrasts on two endpoints (Section II-D). Qwen2.5-VL-72B, sharded across two 80 GB cards, sits at 49.2% and 50.0%. Splitting overlay by whether the red outline was actually drawn (793 against 575 probes) gives 49.3–50.7% with it and 48.7–51.3% without




Fig. 6. **Visual information decay in 2D.** Accuracy on image-required probes as blur σ grows (dashed with markers) and the no-image floor ($\times$); L^* is the smallest σ at which residual visual contribution falls below 20% of the clean-image contribution.

TABLE IV

VOLUMETRIC HEADLINE ON THE GROWTH-MATCHED SUBSET (1,368 PROBES, LABELS EXACTLY BALANCED, CHANCE 50%). GATE: KNOWN-ANSWER CONTROLS PASSED OF 140. MODAL: SHARE OF ANSWERS ON THE MOST FREQUENT OPTION. PAIR: MATCHED PAIRS ANSWERED IDENTICALLY, WHICH GEOMETRY FORBIDS.

Model	Input	Gate	Acc (95% CI)	Modal	Pair
SmolVLM2-2.2B	montage	121	50.0 [48.2, 51.8]	100.0	99.8
Qwen2.5-VL-3B	montage	140	50.0 [48.2, 51.8]	99.6	99.4
Qwen2.5-VL-7B	montage	140	50.0 [48.2, 51.8]	100.0	100.0
Qwen2.5-VL-32B	montage	140	50.6 [48.6, 52.6]	51.5	86.7
Qwen3-VL-8B	montage	140	48.7 [46.6, 50.7]	79.5	88.8
InternVL3-8B	montage	140	50.1 [48.0, 52.2]	72.6	79.4
InternVL3-14B	montage	140	50.1 [48.4, 51.9]	96.5	94.1
LLaVA-OneVision-7B	montage	140	49.5 [47.7, 51.3]	99.0	99.8
Idefics3-8B	montage	82	50.0 [48.2, 51.8]	100.0	100.0
Pixtral-12B	montage	140	51.0 [48.8, 53.3]	61.3	86.6
M3D-LaMed-Phi3-4B	native	107	49.9 [48.1, 51.7]	80.1	93.9
M3D-LaMed-Llama2-7B	native	80	50.1 [48.2, 52.0]	80.5	93.7
Med3DVLM-7B	native	80	50.0 [48.2, 51.8]	100.0	100.0

TABLE V

VOLUMETRIC AUDIT UNDER THE 2D COLUMNS, GROWTH-MATCHED SUBSET. PR, NEG AND SDR ARE THE TEXT-SIDE PERTURBATIONS ON THE PUBLISHED RENDERING; LPA AND SFR ARE CONTENT-FREE CALIBRATED, SFR ON THE CLEANEST TRAP FAMILY; VGR IS POOLED OVER THE FOUR ORGANS; CS FOLLOWS EQ. (3) WITH Safe = 100 - SFR. — NOT RUN.

Model	Original	PR	NEG	SDR	LPA	SFR ↓	VGR	CS
Qwen2.5-VL-7B	50.0	50.0	50.0	50.0	64.5	43.7	+0.0	51.9
InternVL3-8B	50.1	49.2	49.9	49.4	—	71.5	-6.6	40.2
Qwen3-VL-8B	48.7	47.8	50.0	49.7	—	—	-1.0	—
Qwen2.5-VL-32B	50.6	51.4	49.5	49.8	—	73.2	-5.4	38.9
M3D-LaMed-Phi3-4B	49.9	—	—	—	55.0	96.4	-1.2	—

it. The share of matched pairs answered identically rises with the information supplied for three of the four non-degenerate models: InternVL3-8B 81→91%, Qwen2.5-VL-

TABLE VI

IDENTIFICATION CONTROL ON THE GROWTH-MATCHED SUBSET. PLAIN IS THE PUBLISHED CONDITION; IDENTIFIED GIVES JOINT-VISIBILITY SLICES, BOTH OUTLINES AND A 10 MM SCALE BAR. POOLED ACCURACY; INTERVALS FOR THE TWO CELLS THAT EXCLUDE 50 ARE GIVEN IN THE TEXT.

Model	plain	bestslice	overlay	identified
Qwen2.5-VL-7B	50.0	50.0	50.0	50.0
InternVL3-8B	50.0	50.3	49.6	50.4
Qwen3-VL-8B	48.2	49.9	49.5	52.4
Qwen2.5-VL-32B	50.7	51.2	51.0	52.6
Qwen2.5-VL-72B	49.2	—	—	50.0

32B 90→92%, Qwen2.5-VL-72B 89→99%. Qwen2.5-VL-7B answers “no” to all 1,368 probes under all four conditions.

5) *Sub-task decomposition*: The four sub-tasks are asked separately on the identified rendering, in the same four-option format and under the same scoring (300 probes per cell, stratified over organs, chance 25%). The four models with clean response channels localise the red-outlined lesion to its organ at 67.7%, 41.0%, 81.7% and 76.0% (Qwen2.5-VL-7B, InternVL3-8B, Qwen3-VL-8B, Qwen2.5-VL-32B). On an unannotated rendering with the legend dropped, the same question scores 57.3%, 15.0%, 49.7% and 52.0%. They name the cyan-outlined target at 34.3%, 31.3%, 60.0% and 45.0%. They estimate the separation in millimetres, with a 10 mm bar in the panel, at 25.7%, 26.3%, 23.0% and 28.3% (Fig. 7); intervals are [20.7, 31.0], [21.6, 31.3], [18.3, 27.9] and [23.1, 33.8]. Answers collapse onto one or two buckets (Qwen2.5-VL-32B answers “15” on 298 of 300 items), while the correct buckets are near-uniform (85/67/79/69). When the distance question is asked under richer input, InternVL3-8B scores 27.7%, 26.0% and 25.3% on three planes, twenty-five contiguous axial slices and magnified crops. Qwen2.5-VL-32B scores 28.3% on all three and answers “15” on 300 of 300 items under the last two. All eight intervals contain 25. InternVL3-8B’s modal share rises from 48% to 76% under twenty-five slices.

6) *Metric-channel control*: The distance sub-task is at chance on every background (Fig. 8). With two synthetic outlines on uniform grey panels and a 10 mm bar, the four models score 25.7%, 26.7%, 25.7% and 28.3% (Qwen2.5-VL-7B, InternVL3-8B, Qwen3-VL-8B, Qwen2.5-VL-32B); with the same synthetic outlines drawn over the real anatomy, 25.3%, 26.0%, 25.0% and 28.3%; on the real lesion and target, 26.3%, 27.7%, 23.7% and 28.3%. All twelve volume-clustered intervals contain 25%, and each model keeps its answer habit across backgrounds: Qwen2.5-VL-32B answers “15” on 296–300 of 300 items whether the outlines enclose anatomy or nothing at all. The metric channel is absent on a blank background as it is on a CT, so the deficit belongs to the models and the clinical task inherits it. The answer given in its place is a habit: each model’s modal answer and its share are the same on all three backgrounds (Fig. 8c).

7) *What the decision variable tracks*: Two manipulations separate what the decision variable responds to (Fig. 9a). Within a matched pair, the volume, rendering and target are identical and only the growth number in the question differs, with opposite labels. The variable ranks the two members




Fig. 7. **Perception works, naming partly works, metric estimation is absent.** (a) Sub-task accuracy on the identified rendering with volume-clustered 95% intervals: which organ holds the red-outlined lesion, the same with outlines and legend removed, which structure is outlined in cyan, and how many millimetres separate them. (b) Where the 300 distance answers land against the near-uniform gold: every model collapses onto one or two buckets, Qwen2.5-VL-32B onto “15” for 298 items.




Fig. 8. **Two ways the evidence contract breaks in the volume: a channel the models lack on any image, and a clinical sentence that disables the arithmetic they have.** (a) One item’s axial panel under three backgrounds: synthetic outlines on uniform grey, the same synthetic outlines drawn over the real anatomy after its annotations were erased, and the real lesion and target. (b) Distance sub-task accuracy per model and background with volume-clustered 95% intervals; the synthetic gaps follow the real gap distribution and each gold is recomputed from the drawn masks. (c) The modal answer and its share on each background: the answer habit does not change with what the outlines enclose. (d) The comparison the probe reduces to, posed three ways on the 1,368 growth-matched probes: as two numbers, as the same two numbers inside the clinical sentence with no image, and as the published image.

correctly at AUROC 0.61–0.96 for nine models (Qwen2.5-VL-72B 0.959, Qwen3-VL-8B 0.891, Idefics3-8B 0.879, InternVL3-8B 0.855, Qwen2.5-VL-32B 0.853, InternVL3-14B 0.787, Qwen2.5-VL-7B 0.775, Med3DVLM-7B 0.737, LLaVA-OneVision-7B 0.610), so the number is read and moves the score in the right direction. Two models rank it backwards (SmolVLM2-2.2B 0.39, M3D-LaMed-Llama2-7B 0.42). Across volumes that share an identical question sentence and differ only in the image, the same variable ranks the label at 0.46–0.53 for every model, and every interval contains 0.5. The variable therefore depends on the sentence and not on the volume. Scale sharpens the reading of the number without changing the label. Across Qwen2.5-VL 3B, 7B, 32B and 72B, the within-pair AUROC rises from 0.54 to 0.78, 0.85 and 0.96, while the label AUROC stays at 0.50, 0.48, 0.51 and 0.50 (Fig. 9b). Full-corpus accuracy rises from 50.0% to 56.0% between 3B and 32B on the growth cue alone, while growth-matched accuracy stays at 50.0–50.6%.

The same picture holds at the level of answers. Over paired sighted and blind runs on identical probes we measure, per model, the mean absolute perturbation of the option log-probabilities, the mean decision gap between the top two options in the sighted arm, and the share of probes whose argmax differs between arms. Six models perturb their scores measurably and change almost no answers. Qwen2.5-VL-7B moves 0.355 nats against a 1.657-nat decision gap and flips 0.0% of 2,262 probes [0.0, 0.0], Idefics3-8B and Med3DVLM-7B flip 0.0%, and SmolVLM2, Qwen2.5-VL-3B and LLaVA-OneVision flip 0.1–1.2%. Five models show the opposite pattern. InternVL3-14B changes its answer on 92.0% [90.7, 93.4] of probes when shown the volume, Qwen3-VL-8B on 77.4%, Pixtral-12B on 57.3%, Qwen2.5-VL-32B on 48.9% and InternVL3-8B on 30.3%, while their confound-free accuracies are 50.1%, 48.7%, 51.0%, 50.6% and 50.1%. The two native M3D-LaMed variants lie between the two groups, at 18.5% and 19.7%.

8) *Ceilings*: The reference rule applied to stored prove-nance scores 100.0% on 8,476 probes. All five models run on the ceiling arms (Qwen2.5-VL 7B, 32B and 72B, InternVL3-8B and Qwen3-VL-8B) score 100.0% on the numeric oracle, “Is 21.4 greater than or equal to 18.16?”, with balanced predictions and no pair answered identically; the four with a greedy run score 100.0% under greedy decoding as well. The text oracle supplies the same two numbers inside the clinical sentence and removes the image: Qwen2.5-VL-7B, InternVL3-8B and Qwen3-VL-8B score 50.0% with one constant answer on 1,368 of 1,368 items; Qwen2.5-VL-32B 51.1% (“no” on 1,353); Qwen2.5-VL-72B 52.4% (“yes” on 1,335). Greedy decoding yields the same constants for four of the five models; Qwen2.5-VL-32B leaves 82% of its greedy answers unparseable and scores 18.0%. The three forms of one comparison sit side by side in Fig. 8d: the numbers alone are solved, the numbers inside the clinical sentence are not, and the image adds nothing to the sentence.

9) *Grounding*: The visual grounding ratio is negative in every organ for both models with a verified response channel, −4.5 to −6.3 for Qwen2.5-VL-32B and −5.7 to −8.3 for InternVL3-8B, and it is +0.0 in all four organs for Qwen2.5-

TABLE VII

THE SCORING RULE DECIDES THE COMPOSITE. FIVE TRAP FAMILIES, 3,200 PROBES, IDENTICAL FORWARD PASSES; ONLY THE RULE DIFFERS. RAW = LIKELIHOOD OVER OPTION STRINGS; CALIB. = CONTENT-FREE CALIBRATION AGAINST A COMMON BASELINE.

Metric	M3D-LaMed-Phi3		Qwen2.5-VL-7B	
	Raw	Calib.	Raw	Calib.
Acc _{orig}	0.0	45.5	0.5	39.5
Trap refused	100.0	2.0	100.0	76.7
Refusal on answerable	85.8	2.8	97.9	17.5
SFR _w ↓	0.0	97.3	0.0	23.3
CS	0.0	7.1	2.0	49.1

VL-7B. A four-arm decomposition on the same items adds the untouched volume (*full*) and a blank one (*zero*) to the two ROI arms (Fig. 10). Removing the evidence region costs −3.5 to +1.8 points (*full* minus *roi_masked*); removing everything but the evidence region costs +2.2 to +8.1 (*full* minus *roi_only*); the whole volume is worth +4.3 to +10.4 over blank (*full* minus *zero*). For lung and colon, a second run that records the mask fill per item gives the same values to the decimal.

10) *Safety and the scoring rule*: Table VII reports five trap families over 3,200 probes for two models under identical forward passes scored two ways. Under raw likelihood both models refuse 100% of traps and score SFR_w = 0.0%, while also selecting the refusal string on 85.8% and 97.9% of answerable probes; their composites are 0.0 and 2.0. Under content-free calibration the same passes separate: Qwen2.5-VL-7B declines 76.7% of unsatisfiable premises, evenly across tiers, at SFR_w = 23.3% and CS 49.1; M3D-LaMed declines 2.0% at SFR_w = 97.3% and CS 7.1. Raw scoring reports SFR_w = 0.0% for every model on every family (four models, two input routes, eight arms). On the cleanest trap family, calibrated SFR is 43.7% (Qwen2.5-VL-7B), 71.5% (InternVL3-8B), 73.2% (Qwen2.5-VL-32B) and 96.4% (M3D-LaMed).

11) *Pretraining overlap*: A probe asking which public collection a scan comes from carries its own positive control, since naming the organ can be answered by looking. Over 80 volumes per cell and four models, the control scores 50.0–57.5% while dataset identification scores 0.0–5.0% and case identification 17.5–26.2%. The modal dataset answer is a constant. Three of four models place 58–76 of 80 answers on “the TCIA Pancreas-CT collection”, from which none of these volumes come.

C. Cross-Modality Comparison

The two audits share a framework, metric set and composite, so the differences below are differences in what the modality affords.

1) *Image contribution*: In 2D the image accounts for the score. Withholding it from six models lowers original-probe accuracy by 24.7 to 69.7 points (Table VIII), two text-only baselines score 1.7% and 4.0% on the same probes, and under the blur sweep the clean-image contribution is 34.7–56.3 points, with language taking over only at $L^* = 16$ –64 px




Fig. 9. **The decision variable reads the question, not the volume.** (a) AUROC of $\log p_{\text{yes}} - \log p_{\text{no}}$ on the growth-matched subset with volume-clustered 95% intervals, three comparisons per model: within the 363 complete matched pairs, where volume and rendering are identical and only the growth number in the question differs; against the computed label over all 1,368 probes; and across the 182 probes whose question sentence is identical and only the volume differs. [†]Native volumetric input. (b) The growth-number and label AUROCs across Qwen2.5-VL 3B–72B on the published rendering.




Fig. 10. **Removing the evidence costs nothing; removing everything else costs the volume and collapses the answer.** (a, b) Accuracy under the untouched volume, the evidence region removed, and only the evidence region kept; a blank volume scores 50.0 (dashed). (c, d) Share of answers on the modal option under each arm: with only the evidence region kept the models fall to a near-constant answer (labelled), which is why that arm sits at chance.

(Fig. 6). In 3D the same subtraction on the growth-matched subset runs from -1.3 to $+1.0$ points across thirteen models. The blur sweep on two of these models gives $L^* \in \{0, 8\}$ with clean-image contributions of 2.8 and 0.0 points, and the decision variable ranks the label at 0.47–0.52 with the volume present. Under one protocol, the image is decisive in one regime and inert in the other.

2) *Text robustness*: Paraphrase, negation and specificity drop span 29.7–80.7%, 15.6–83.1% and 25.0–79.9% across

the sixteen 2D models, with the radiologist at 92.7, 90.7 and 91.5; in 3D the four clean-channel models sit at 47.8–51.4% on all three, and 99.2–100% of negated probes receive “no”.

3) *Grounding*: VGR spans $+3.6$ to $+49.5$ pp across capable 2D models and is positive for every one of them; in 3D it spans -8.3 to $+0.8$ pp, is negative in every organ for both verified-channel models, and is exactly $+0.0$ in four of four organs for one model. The radiologist’s 2D VGR is $+4.5$.

TABLE VIII
IMAGE CONTRIBUTION ACROSS MODALITIES. 2D: ORIGINAL-PROBE ACCURACY WITH AND WITHOUT THE IMAGE ON THE SAME PROBES, AND TWO TEXT-ONLY BASELINES. 3D: SIGHTED AND BLIND ACCURACY ON THE GROWTH-MATCHED SUBSET.

	Model	Image	No image	Δ
2D	GPT-5.5	75.0	5.3	+69.7
	Claude Opus 4.7	78.7	16.0	+62.7
	GPT-4o	59.3	6.3	+53.0
	Gemini 3.1 Flash-Lite	73.0	22.7	+50.3
	Gemini 3 Flash	77.0	36.0	+41.0
	GPT-5.4-nano	31.7	7.0	+24.7
	DeepSeek-V4-Pro (text only)	—	1.7	—
	DeepSeek-V4-Flash (text only)	—	4.0	—
	Qwen3-VL-8B	48.7	50.0	−1.3
3D	InternVL3-8B	50.1	50.0	+0.1
	Qwen2.5-VL-32B	50.6	49.9	+0.7
	Pixtral-12B	51.0	50.0	+1.0
	all thirteen (range)	48.7–51.0	49.5–50.0	[−1.3, +1.0]

4) *Safety*: 2D SFR ranges 22.0–62.8% against the radiologist’s 5.8%. Calibrated volumetric SFR on the cleanest trap family is 43.7%, 71.5%, 73.2% and 96.4%. The first value lies inside the 2D range and the other three above it, and the highest belongs to the only purpose-built medical volumetric model in the set. Uncalibrated, every volumetric model scores $\text{SFR}_w = 0.0\%$.

IV. DISCUSSION

In this paper we evaluate evidence-conditional trustworthiness of medical VLMs in two regimes, a clinician-annotated 2D suite and a volumetric suite whose labels are computed from geometry, and audit twenty-nine model configurations under one framework. The 2D finding survives the change of modality and worsens, the volumetric labels cost nothing to produce and cannot have been memorised, and the composite probe decomposes into steps that fail selectively. We discuss which models carry the volumetric claim, why one chance-level accuracy hides two different failures, what the volumetric regime adds to the framework, and what follows for clinical use.

A. Which Models Carry the Volumetric Claim

Thirteen models at 48.7–51.0% are not thirteen independent measurements, and the modal-share bound identifies which of them measure anything. On a label-balanced binary corpus, a model that places a fraction m of its answers on one option has accuracy confined to $[m - \frac{1}{2}, \frac{3}{2} - m]$ regardless of what it computes. Seven of the thirteen models sit at $m \geq 0.96$, so their accuracy near 50% follows from their answer distribution, and InternVL3-14B’s 50.1% is confined to [46.5, 53.5] before any reasoning is considered. Removing those seven and the three native systems that fail the response-channel gate leaves four models that carry the volumetric accuracy claim, Qwen3-VL-8B, InternVL3-8B, Pixtral-12B and Qwen2.5-VL-32B, each with an interval containing 50. For this reason the modal share is reported beside every accuracy in Table IV. The same bound applies to the sub-task result. Qwen2.5-VL-32B answers “15 mm” on 298–300 of 300 items in every arm, so

its 28.3% is the base rate of that bucket. The distance finding therefore rests on InternVL3-8B, which spreads its answers (48–50% modal share) and still scores 26.3%.

B. Two Failures Behind One Number

Chance-level accuracy is compatible with two states that accuracy alone cannot separate. Either nothing usable reaches the decision, or something does and is out-voted. The paired sighted and blind runs separate the two (Section III-B). Six models are in the first state, since the volume perturbs their option scores by less than the decision gap and changes almost no answers. Five models are in the second state, which is harder to explain away. InternVL3-14B changes 92% of its answers when shown the volume and remains at 50.1%, so it is not ignoring the image. The volume reaches the decision and changes most of it without making it right, and the material it changes it with is the question. The same decision variable ranks the growth number in the sentence at AUROC up to 0.96 and the label at 0.47–0.52, and larger Qwen models read the number better while the label does not move (Fig. 9). The distinction matters for remedies. Amplifying an under-weighted visual signal is the right instrument for the first group, but it cannot help the second, where nothing is under-weighted and the content of what arrives is wrong.

The sub-task decomposition identifies that content. With both structures outlined and a metric scale bar in the panel, the same models that localise the lesion at up to 81.7% and answer *is 21.4 greater than or equal to 18.16?* at 100% estimate the separation between the two outlines at chance. Categorical information survives the pipeline, and metric information does not. Four ordinary explanations were tested directly. The input was rebuilt to give the model what a reader would be given, the visible volume was multiplied by $5.7\times$ and the structures were magnified, a tenfold parameter range was spanned, and the anatomy was replaced by two synthetic outlines on a blank background. None of these changed the metric result. The missing metric channel is therefore a property of the models and not of CT: they cannot read a distance off a blank page with a scale bar, which agrees with reports that general-domain VLMs fail elementary geometric readouts [40], [41] and that metric spatial reasoning has required dedicated training [42]. The clinical task inherits that limit. The failure is a missing metric channel in models that otherwise see what they are shown, and remedies aimed at the encoder are aimed at the part that already works.

C. A Component That Is Not Visual

Supplying the two distances in words and removing the image still collapses every model to a single constant answer, while the same comparison stripped of clinical wording is solved at 100%. Greedy decoding gives the same constants for four of the five models and an unparseable answer for most items of the fifth, so this is not an artefact of likelihood scoring. The task is solvable, as the ceiling arms show, and the collapse reproduces with no image present. Any account of the volumetric null must therefore include a component that is not visual. The clinical sentence costs the model the arithmetic

it otherwise performs. This component was measurable only because the labels were computed, which allowed the same items to be asked again in text with their gold intact.

D. What the Volumetric Regime Adds to the Framework

The volumetric regime required one control that we expect to generalise. The *response-channel control* was unnecessary in 2D, where every audited model answers seven known-answer questions correctly, and it is indispensable in 3D, where none of the three native volumetric systems passes. A pinned output channel scores 50% and cannot be told apart in an accuracy table from systems that fail for entirely different reasons, and the four models that fail here are not the four an architectural reading would predict. The control also anticipates the safety axis. Every failure is a failure of polarity, near-perfect in one answer direction and near-chance or worse in the other. Among the four models measured on both instruments, the one that answers the controls’ “no” items at 61% fabricates on 96.4% of calibrated traps, while the three that answer them at 100% fabricate on 43.7–73.2%. Seven questions on twenty volumes anticipate what the trap families take days to establish. A model that cannot produce a negative answer about the image does not produce one about the anatomy either.

Two instruments change meaning in the volumetric regime. The grounding contrast subtracts an arm that removes the evidence region from one that removes everything else. In 2D the two arms occupy comparable image area and the difference approximates grounding. In a volume the region is roughly 1% of the voxels and its complement 99%, so the same difference confounds grounding with intactness, and VGR inverts instead of shrinking (Fig. 10). A Ground axis at its floor still contributes about 50 to a harmonic mean and reads as partial competence. We recommend that in the volumetric regime the Ground axis gate the composite instead of averaging into it, so that a system whose matched-pair violation exceeds a threshold has its CS reported as undefined. The second instrument is the scoring rule. The 2D suite is scored by an emitted letter, so the refusal-string likelihood artefact of Table VII does not arise there in the same form. However, any sighted-versus-blind difference scored against per-arm baselines inherits the second artefact we found, namely that a model’s content-free answer differs with and without an image, and we recommend a common baseline.

E. Clinical Reading and Deployment

The property that fails is the one a distance-to-structure workflow depends on. A system that perceives an outlined lesion and names the adjacent structure, but cannot say whether they are 5 or 35 mm apart, answers a contact question fluently and at chance, and nothing in its output marks the difference. The only purpose-built medical volumetric model in the set fabricates on 96.4% of calibrated traps, and none of the three native systems can answer “*Is this a CT scan?*” reliably. This is the volumetric form of the 2D finding that medical fine-tuning does not yet translate into safer behaviour. The corpus bounds the reading in one direction. Probe generation

excludes every lesion within 2 mm of its target, so the cases in which management actually turns are outside this audit by construction.

Two consequences follow for use. An audit should report clean-input capability, the trap silent-failure rate and its risk-weighted form, the grounding decomposition, language-prior controls and the response-channel gate as separate axes, because they fail independently here and a composite hides which one failed. The safest mitigation sits upstream of the model. A clinical pipeline should validate image presence, integrity, modality and resolution before a request reaches a VLM, and it should route a failed check to a reader instead of to a system that answers fluently through it. A model that cannot say that a scan is not a CT will not say that a lesion does not reach the aorta.

V. CONCLUSION

We presented a two-regime evaluation of evidence-conditional trustworthiness in medical VLMs. The 2D suite, authored by four radiologists, showed that capability and safe failure diverge and that silent failure concentrates where harm is greatest. Replacing annotation with computation in CT removed the two objections that suite could not answer from inside, scale and pretraining overlap, at no annotation cost, and the finding replicated and worsened. Decomposing the computed probe localised the volumetric failure to a single absent step, metric estimation, which fails on synthetic outlines as it fails on anatomy; the decision variable was shown to track the question rather than the volume, and the audit exposed two instruments that do not survive the change of modality.

Future work includes a volumetric reader study on the released 104-probe form, probes on lesions within 2 mm of their target, inference-time compute in the volumetric regime, and the extension of the computed construction to further geometric relations and modalities. Both suites, the harness and every cached model output are released, so the computed construction is available as a low-cost test of whether a medical VLM knows when its evidence has failed.

REFERENCES

- [1] C. Li *et al.*, “LLaVA-Med: Training a large language-and-vision assistant for biomedicine in one day,” in *Adv. Neural Inf. Process. Syst. (NeurIPS), Datasets and Benchmarks Track*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.00890>
- [2] J. Chen *et al.*, “Towards injecting medical visual knowledge into multimodal LLMs at scale,” in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2024, pp. 7346–7370. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.418/>
- [3] OpenAI, “GPT-4o system card,” 2024. [Online]. Available: <https://openai.com/index/gpt-4o-system-card/>
- [4] Anthropic, “The Claude 3 model family: Opus, sonnet, haiku,” 2024. [Online]. Available: <https://www.anthropic.com/news/claude-3-family>
- [5] Gemini Team Google, “Gemini: A family of highly capable multimodal models,” *arXiv preprint arXiv:2312.11805*, 2023. [Online]. Available: <https://arxiv.org/abs/2312.11805>
- [6] M. Moor *et al.*, “Med-Flamingo: A multimodal medical few-shot learner,” in *Proceedings of the 3rd Machine Learning for Health Symposium*, 2023, pp. 353–367. [Online]. Available: <https://proceedings.mlr.press/v225/moor23a.html>
- [7] C. Wu, X. Zhang, Y. Zhang, Y. Wang, and W. Xie, “Towards generalist foundation model for radiology by leveraging web-scale 2d and 3d medical data,” *arXiv preprint arXiv:2308.02463*, 2023. [Online]. Available: <https://arxiv.org/abs/2308.02463>

- [8] M. Asadi *et al.*, “Mirage: The illusion of visual understanding,” *arXiv preprint arXiv:2603.21687*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.21687>
- [9] L. Chen *et al.*, “BabyVision: Visual reasoning beyond language,” *arXiv preprint arXiv:2601.06521*, 2026. [Online]. Available: <https://arxiv.org/abs/2601.06521>
- [10] Y. Goyal, T. Khot, D. Summers-Stay, D. Batra, and D. Parikh, “Making the V in VQA matter: Elevating the role of image understanding in visual question answering,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 6904–6913. [Online]. Available: https://openaccess.thecvf.com/content_cvpr_2017/html/Goyal_Making_the_v_CVPR_2017_paper.html
- [11] V. Agarwal, R. Shetty, and M. Fritz, “Towards causal VQA: Revealing and reducing spurious correlations by invariant and covariant semantic editing,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2020, pp. 9690–9698. [Online]. Available: https://openaccess.thecvf.com/content_CVPR_2020/html/Agarwal_Towards_Causal_VQA_Revealing_and_Reducing_Spurious_Correlations_by_Invariant_CVPR_2020_paper.html
- [12] J. J. Lau, S. Gayen, A. Ben Abacha, and D. Demner-Fushman, “A dataset of clinically generated visual questions and answers about radiology images,” *Sci. Data*, vol. 5, p. 180251, 2018. [Online]. Available: <https://www.nature.com/articles/sdata2018251>
- [13] B. Liu, L.-M. Zhan, L. Xu, L. Ma, Y. Yang, and X.-M. Wu, “SLAKE: A semantically-labeled knowledge-enhanced dataset for medical visual question answering,” in *Proc. IEEE Int. Symp. Biomed. Imag. (ISBI)*, 2021, pp. 1650–1654. [Online]. Available: <https://arxiv.org/abs/2102.09542>
- [14] X. Zhang *et al.*, “PMC-VQA: Visual instruction tuning for medical visual question answering,” *arXiv preprint arXiv:2305.10415*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.10415>
- [15] X. Chen *et al.*, “Auditing frontier vision-language models for trustworthy medical VQA: Grounding failures, format collapse, and domain adaptation,” *arXiv preprint arXiv:2604.27720*, 2026. [Online]. Available: <https://arxiv.org/abs/2604.27720>
- [16] A. Rohrbach, L. A. Hendricks, K. Burns, T. Darrell, and K. Saenko, “Object hallucination in image captioning,” in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018, pp. 4035–4045. [Online]. Available: <https://aclanthology.org/D18-1437/>
- [17] Y. Li, Y. Du, K. Zhou, J. Wang, X. Zhao, and J.-R. Wen, “Evaluating object hallucination in large vision-language models,” in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2023, pp. 292–305. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.20/>
- [18] J. Chen *et al.*, “Detecting and evaluating medical hallucinations in large vision language models,” *arXiv preprint arXiv:2406.10185v1*, 2024. [Online]. Available: <https://arxiv.org/abs/2406.10185v1>
- [19] A. Chang, L. Huang, P. Bhatia, T. Kass-Hout, F. Ma, and C. Xiao, “MedHEval: Benchmarking hallucinations and mitigation strategies in medical large vision-language models,” *arXiv preprint arXiv:2503.02157*, 2025. [Online]. Available: <https://arxiv.org/abs/2503.02157>
- [20] Z. Gu, C. Yin, F. Liu, and P. Zhang, “MedVH: Towards systematic evaluation of hallucination for large vision language models in the medical context,” *arXiv preprint arXiv:2407.02730*, 2024. [Online]. Available: <https://arxiv.org/abs/2407.02730>
- [21] A. Heiman, X. Zhang, E. Chen, S. E. Kim, and P. Rajpurkar, “FactCheXcker: Mitigating measurement hallucinations in chest x-ray report generation models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2025, pp. 30 787–30 796.
- [22] D. Nguyen *et al.*, “Localizing before answering: A benchmark for grounded medical visual question answering,” in *Proc. Int. Joint Conf. Artif. Intell. (IJCAI)*, 2025, pp. 7670–7678. [Online]. Available: <https://www.ijcai.org/proceedings/2025/853>
- [23] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” in *Advances in Neural Information Processing Systems*, 2017. [Online]. Available: <https://papers.nips.cc/paper/7073-selective-classification-for-deep-neural-networks>
- [24] J. Ren, Y. Zhao, T. Vu, P. J. Liu, and B. Lakshminarayanan, “Self-evaluation improves selective generation in large language models,” *arXiv preprint arXiv:2312.09300*, 2023. [Online]. Available: <https://arxiv.org/abs/2312.09300>
- [25] B. Wen *et al.*, “Know your limits: A survey of abstention in large language models,” *Transactions of the Association for Computational Linguistics*, vol. 12, pp. 1412–1430, 2024. [Online]. Available: <https://arxiv.org/abs/2407.18418>
- [26] P. Kirichenko, M. Ibrahim, K. Chaudhuri, and S. J. Bell, “AbstentionBench: Reasoning LLMs fail on unanswerable questions,” *arXiv preprint arXiv:2506.09038*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.09038>
- [27] N. Madhusudhan, V. Yadav, and A. Lacoste, “Knowing when not to answer: Evaluating abstention in multimodal reasoning systems,” *arXiv preprint arXiv:2604.14799*, 2026. [Online]. Available: <https://arxiv.org/abs/2604.14799>
- [28] N. Moratelli, C. Davis, L. F. R. Ribeiro, B. Byrne, and G. Iglesias, “Benchmarking deflection and hallucination in large vision-language models,” *arXiv preprint arXiv:2604.12033*, 2026. [Online]. Available: <https://arxiv.org/abs/2604.12033>
- [29] J. Li, L. Fang, H. Zhao, S. Duan, and G. Liu, “Judge before answer: Can MLLM discern the false premise in question?” *arXiv preprint arXiv:2510.10965*, 2025. [Online]. Available: <https://arxiv.org/abs/2510.10965>
- [30] S. Wang *et al.*, “A novel evaluation benchmark for medical LLMs: Illuminating safety and effectiveness in clinical domains,” *arXiv preprint arXiv:2507.23486*, 2025. [Online]. Available: <https://arxiv.org/abs/2507.23486>
- [31] K. Singhal *et al.*, “Large language models encode clinical knowledge,” *Nature*, vol. 620, pp. 172–180, 2023. [Online]. Available: <https://www.nature.com/articles/s41586-023-06291-2>
- [32] M. Antonelli *et al.*, “The Medical Segmentation Decathlon,” *Nat. Commun.*, vol. 13, no. 1, p. 4128, 2022.
- [33] O. Pelka, S. Koitka, J. Rückert, F. Nensa, and C. M. Friedrich, “Radiology objects in COntext (ROCO): A multimodal image dataset,” in *Proc. MICCAI Workshop Large-Scale Annotation Biomed. Data Expert Label Synth. (LABELS)*, 2018, pp. 180–189.
- [34] A. E. W. Johnson *et al.*, “MIMIC-CXR, a de-identified publicly available database of chest radiographs with free-text reports,” *Sci. Data*, vol. 6, p. 317, 2019. [Online]. Available: <https://www.nature.com/articles/s41597-019-0322-0>
- [35] J. Irvin *et al.*, “CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison,” in *Proc. AAAI Conf. Artif. Intell.*, vol. 33, no. 1, 2019, pp. 590–597.
- [36] J. Wasserthal *et al.*, “TotalSegmentator: Robust segmentation of 104 anatomic structures in CT images,” *Radiol. Artif. Intell.*, vol. 5, no. 5, p. e230024, 2023.
- [37] C. R. Maurer, R. Qi, and V. Raghavan, “A linear time algorithm for computing exact Euclidean distance transforms of binary images in arbitrary dimensions,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 25, no. 2, pp. 265–270, 2003.
- [38] A. Holtzman, P. West, V. Shwartz, Y. Choi, and L. Zettlemoyer, “Surface form competition: Why the highest probability answer isn’t always right,” in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2021, pp. 7038–7051.
- [39] T. Z. Zhao, E. Wallace, S. Feng, D. Klein, and S. Singh, “Calibrate before use: Improving few-shot performance of language models,” in *Proc. Int. Conf. Mach. Learn. (ICML)*, 2021, pp. 12 697–12 706.
- [40] P. Rahmzadehgervi, L. Bolton, M. R. Taesiri, and A. T. Nguyen, “Vision language models are blind,” *arXiv preprint arXiv:2407.06581*, 2024.
- [41] A. Kamath, J. Hessel, and K.-W. Chang, “What’s “up” with vision-language models? Investigating their struggle with spatial reasoning,” *arXiv preprint arXiv:2310.19785*, 2023.
- [42] B. Chen *et al.*, “SpatialVLM: Endowing vision-language models with spatial reasoning capabilities,” *arXiv preprint arXiv:2401.12168*, 2024.